

# Indexing

RG 8.1, 8.2, 8.3, 8.5

Rasmus Pagh



# Disk crash course

- Relations of large databases are usually stored on hard drives (or SSDs).
- Hard drives can store large amounts of data, but work rather slowly compared to the memory of a modern computer:
  - The time to access a specific piece of data is on the order of  $10^6$  ( $10^4$ ) times slower.
  - The **rate** at which data can be read is on the order of 100 (10) times slower.
- Time for accessing disk may be the main performance bottleneck!



# Parallel data access

- Large database systems (and modern SSDs) use several storage units to:
  - Enable several pieces of data to be fetched in parallel.
  - Increase the total rate of data from disk.
- Systems of several disks are often arranged in so-called RAID systems, with various levels of error resilience.
- Even in systems with this kind of parallelism, the time used for accessing data is often the performance bottleneck.



# Full table scans

When a DBMS sees a query of the form

```
SELECT *  
FROM R  
WHERE <condition>
```

the obvious thing to do is read through the tuples of R and report those tuples that satisfy the condition.

This is called a **full table scan**.



# Selective queries

Consider the query from before

- If we have to report 80% of the tuples in R, it makes sense to do a full table scan.
- On the other hand, if the query is very **selective**, and returns just a small percentage of the tuples, we might hope to do better.



# Point queries

- Consider a selection query with a single equality in the condition:

```
SELECT *  
FROM person  
WHERE birthyear=1975
```

- This is a point query: We look for a single value of "year".
- Point queries are easy if data is sorted by the right attribute.



# Range queries

- Consider a selection query of the form:  
`SELECT *`  
`FROM person`  
`WHERE year(birthdate) BETWEEN`  
`1975 and 1994`
- This is a **range query**: We look for a range of values of "birthdate".
- Range queries are **also** easy if data is sorted by the right attribute.
  - But often not be as selective as point queries.



# Indexes

- To speed up queries the DBMS may build an **index** on the year attribute.
- A database index is similar to an index in the back of a book:
  - For every piece of data you might be interested in (e.g., the attribute value 1975), the index says where to find it.
  - The index itself is organized such that one can quickly do the lookup.
- Looking for information in a relation with the help of an index is called an **index scan**.





# Primary indexes

- If the tuples of a relation are stored sorted according to some attribute, an index on this attribute is called **primary**.
  - Primary indexes make point and range queries on the key *very efficient*.
- Many DBMSs automatically build an index on the primary key of each relation.
  - In MySQL this depends on the storage engine: InnoDB builds an index on the primary key, MyISAM does not.
- A primary index is sometimes referred to as a **clustering** or **sparse** index.



# Secondary indexes

- It is possible to create further indexes on a relation. Typical syntax:

```
CREATE INDEX myIndex ON involved(actorId);
```

- The non-primary indexes are called **secondary** indexes (sometimes **non-clustering** or **dense** indexes)
  - Secondary indexes make *most* point queries on the key more efficient.
  - Secondary indexes make *some* range queries on the key more efficient.



# Multi-attribute indexes

Defining an index on several attributes:

```
CREATE INDEX myIndex  
ON person (name,birthdate);
```

Speeds up point queries such as:

```
SELECT *  
FROM person  
WHERE name='John Doe' and  
birthdate<'1950-1-1'
```

An index on several attributes usually gives index for any **prefix** of these attributes, due to lexicographic sorting.



# Problem session

- What kinds of point and range queries are “easy” when the relation is stored as in the previous example:
  1. A range query on name?
  2. A range query on birthdate?
  3. A point query on birthdate?
  4. A point query on birthdate combined with a range query on firstname?
  5. A point query on firstname combined with a range query on birthdate?



# Index scan vs full table scan

Point and range queries on the attribute(s) of the **primary index** are almost always best performed using an index scan.

**Secondary indexes** should be used with **high selectivity queries**:

As a rule of thumb, a secondary index scan is faster than a full table scan for queries returning less than 10% of a relation.



# Choosing to use an index

- The choice of whether to use an index is made by the DBMS *for every instance of a query*
  - May depend on query parameters
  - Don't have to take indexes into account when writing queries
- Estimating selectivity is done using statistics
  - In MySQL, statistics is gathered by executing statements such as `ANALYZE TABLE` involved



# What speaks against indexing?

- Space usage:
  - Small for primary index
  - Similar to data size for secondary index
- Time usage for keeping indexes updated under data insertion/change:
  - Small to medium for primary index
  - High for secondary index (but has been going down...)



# Other impact of indexes

The DBMS may use indexes in other situations than a simple point or range query.

- Some joins can be executed using a modest number of index lookups
  - May be faster than looking at all data
- Some queries may be executed by only looking at the information in the index
  - **Index only** query execution plan ("covering index").
  - May need to read much less data.





# Index types

Common:

- B-trees (point queries, range queries)
- Hash tables (only point queries, but somewhat faster)
- Bitmap indexes (good for "dense" sets)

More exotic:

- Full text indexes (substring searches)
- Spatial indexes (proximity search, 2D range search,...)
- ... and thousands more



# Conclusion

- Large databases need to be equipped with suitable indexes.
  - Need understanding of what indexes might help a given set of queries.
  - Important distinction: Primary vs secondary.
  - A detailed understanding of various index types is beyond the scope of this course.



# Related course goal

Students should be able to:

- decide if a given index is likely to improve performance for a given query.